ML-Based Feature Importance Estimation for Predicting Unethical Behaviour under Pressure

Pablo Rivas [1]   **Pamela J. Harper** [2]   **John C. Cary** [3]   **William S. Brown** [4]

Abstract

We studied the utility of using machine learning algorithms in the estimation of feature importance and to visualize their dependence on *Ethicality*. Through our analysis and partial dependence plot we found linear relationships among variables and gained insight into features that might cause certain types of ethical behaviour.

1. Introduction

As a result of numerous high-profile ethical lapses by corporations and their employees, research into the contributing factors of ethical conduct has grown. To that end we investigated several specific member attributes and behaviors that impact ethical conduct (Cary et al., 2018; Cary & Rivas, 2017). Our study examined the role of gender, income, and religiosity in shaping ethical conduct, and the degree to which perceptions of pressure might moderate these variables. Using standard statistical analysis such as linear regression and correlation coefficients, we determined that in addition to gender and religiosity, the perception of pressure is a factor in unethical behavior. However, state-of-the-art machine learning (ML) algorithms have also proven to be robust in modeling features in datasets and utilizing intrinsic non-linear transformations over such features to determine the best way to utilize them (López et al., 2018; Rahman et al., 2018; Xiao et al., 2018; Lin et al., 2017; Wei et al., 2016). Thus, this work aims to use ML to determine the relative importance of the feature set in our dataset.

[1]Department of Computer Science, School of Computer Science and Mathematics, Marist College, New York, USA [2]Department of Organization and the Environment, School of Management, Marist College, New York, USA [3]Economics, Accounting, and Finance Department, School of Management, Marist College, New York, USA [4]Management Department, School of Management, Marist College, New York, USA. Correspondence to: Pablo Rivas <Pablo.Rivas@Marist.edu>, Pamela J. Harper <Pamela.Harper@Marist.edu>.

2. Background and Methods

Dataset. In our previous study we administered a questionnaire to 336 subjects in the northeastern United States. The sample group included undergraduate students about to enter the workforce and graduate students who are currently employed. Table 1 in the appendix shows descriptive statistics about the features in the dataset.

Support Vector Machines for Regression (SVRs). If we define a positive constant $C > 0$ describing the trade off between the training error and define a penalizing term on the parameters of an SVR as $||\mathbf{w}||_2^2$ promoting sparser solutions on $\mathbf{w}$. And if we further, let variables ξ_i and ξ_i^* be two sets of nonnegative slack variables that describe an $\epsilon-$insensitive loss function; then we can define an SVR as a predictor over $\mathbf{x}$ with the objective $\min_{\mathbf{w},b,\boldsymbol{\xi},\boldsymbol{\xi}^*} \frac{1}{2}||\mathbf{w}||_2^2 + C \sum_{i=1}^{N} (\xi_i + \xi_i^*)$ with constraints $y_i - \mathbf{w}^T\mathbf{x}_i - b \leq \epsilon + \xi_i$, $\mathbf{w}^T\mathbf{x}_i + b - y_i \leq \epsilon + \xi_i^*$, and $\boldsymbol{\xi}, \boldsymbol{\xi}^* \geq \mathbf{0}$, for $i = 1, 2, \ldots, N$, where $\mathcal{D} = \{\mathbf{x}_i, y_i\}_{i=1}^N$ defines our data set. With this we trained an SVR to model $f(x)$ by successively selecting a single feature from the feature set to record the cross-validated R^2 coefficient using each, where $R^2 = 1 - \frac{\sum(y-f(x))^2}{\sum(y-\bar{y})^2}$. Values of R^2 close to 1 or -1 would indicate high feature importance, while values close to 0 have low predictive value. Figure 1 depicts the results of raking the features using SVRs; this indicates that *Pressure* is one of the best predictors by itself; the rest of the features, individually, are arguably not good predictors. Another feature importance metric we can use with SVRs is in terms of its improvement or worsening of the R^2 coefficient. For this we systematically remove a specific feature and train with the rest to determine the contribution of such feature. First, we establish a baseline coefficient R^* which accounts for training with the full set of features and getting the cross validated score. Then for k-th feature we can quantify its level of contribution by observing the change with respect to the baseline, Δ_{R^*}. The contribution of the k-th feature can be determined as $\Delta_{R_k^*} = |R^*| - |R_k^2|$. Figure 1 depicts the results of our $\Delta_{R_k^*}$ analysis on the feature sets where it can be seen that the removal of the variables *Employment, Education*, and *Pressure* causea a positive $\Delta_{R_k^*}$, i.e., the model significantly drops its predictive ability if these features are removed.
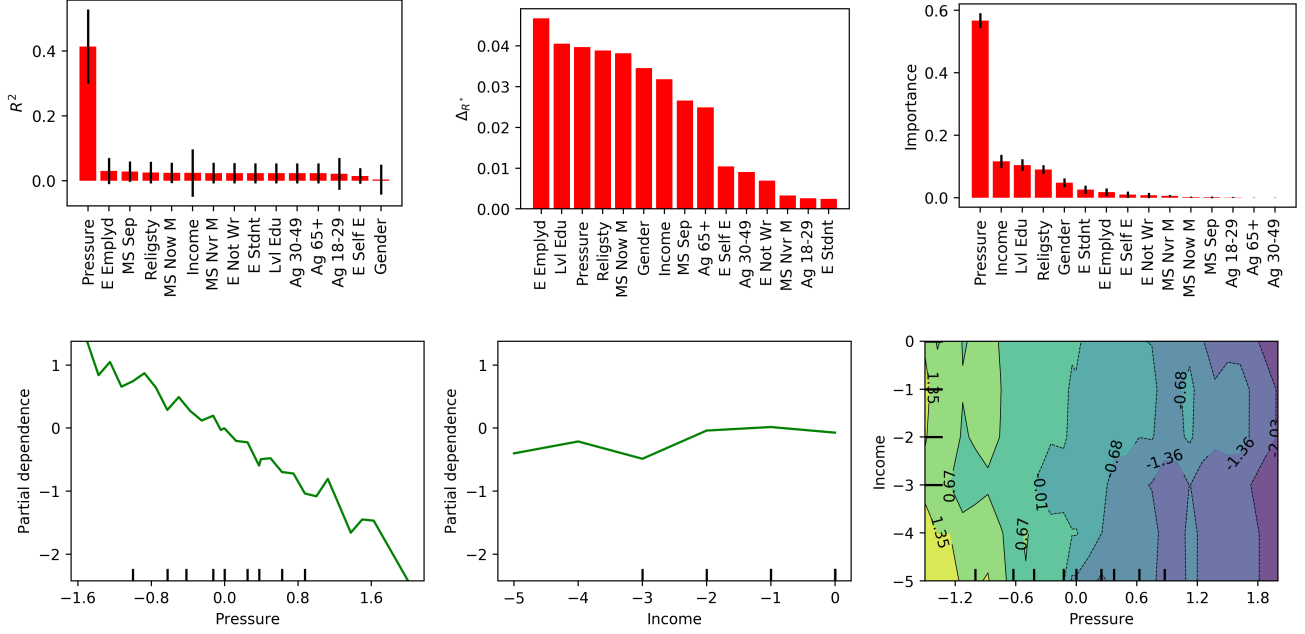
Figure 1. (Top left) R^2 feature importance of isolated features on SVRs. (Top middle) Analysis of $\Delta_{R_k^*}$ on SVRs. (Top right) RFs feature ranking. Partial dependence plots for *Pressure* (bottom left), *Income* (bottom middle), and the combination of both in a color-coded contour plot (bottom right). All SVR experiments used Bayesian optimization to find the best set of hyper-parameters (Louppe, 2017).

Random Forests for Regression (RFs). The theory behind RFs indicates that each tree in the ensemble is constructed using bootstrapping to produce samples and to make usually small trees (Geurts & Louppe, 2011). RFs with M trees and N samples usually have size in the order of $O(MN \log(N))$, and one of the most interesting properties of RFs is that they have high bias and low variance, which made them popular in applications that require stability and automatic feature engineering (Soltaninejad et al., 2017; Pinto et al., 2018). Due to the latter, we used RFs to determine feature importance. Features that are frequently and consistently closer to the root, i.e., that are more *pure*, are considered more important. Figure 1 shows the ranking o the features using RFs, which, consistently with SVRs, show that *Pressure* is highly predictive. *Income*, *Education*, and *Religiosity* seem to have adequate predictive power.

Partial Dependence Analysis. Partial dependence plots have been widely used to visually perceive the importance of features among themselves in order to assess their predictive power over a dependent variable (Lemmens & Croux, 2006). These plots have provided great insight in several areas, from the natural sciences (Isayev et al., 2017), to the legal studies (Berk et al., 2016). We considered *Ethicality* our dependent variable and we chose the most predictive, as found with the earlier ML methods, variables for display and we used a standard Gradient Boosting Regressor

(GBR) as our internal predictor (Peter et al., 2017). Figure 1 shows the partial dependence of *Pressure* (bottom left), the dependence of *Income* (bottom middle), and the interaction of both at the same time (bottom right). Our partial dependence plots depict a linear relationship between *Pressure* and our independent variable; while *Income* shows a quasi-concave relationship that is more evident on the combined plot. Figure 2 (appendix) shows a three-dimensional version of the dependence of both *Pressure* and *Income*.

3. Conclusions

In our previous studies we found that when pressure is introduced into a linear regression model, the *Ethicality* of an individual is easier to predict with high statistical significance (Cary et al., 2018). This paper confirms previous findings by assessing the importance of features using state-of-the-art ML models such as SVRs, RFs, and GBRs. However, until now we are able to visualize the quasi-linear dependence of *Pressure* with our dependent variable, *Ethicality*, and further confirm the quasi-concave dependence behaviour of *Income*. The latter suggests that subjects in both the low end and high end of the income range are predictors of *Ethicality*, while subjects whose income is in the middle range are not predictive on *Ethicality*. Further studies will explore the predictive power of sets of features on each other and not necessarily on the dependent variable *Ethicality*.

References

Berk, R. A., Sorenson, S. B., and Barnes, G. Forecasting domestic violence: A machine learning approach to help inform arraignment decisions. *Journal of Empirical Legal Studies*, 13(1):94–115, 2016.

Cary, J. and Rivas, P. Ethics under pressure? In *Proc. of the 2017 Susilo Symposium, Boston University*, June 2017.

Cary, J. C., Brown, W. S., Harper, P. J., and Rivas, P. Ethics under pressure: A study of the effects of gender, religiosity, and income under the perception of pressure. In *Proceedinfs of the 25th International Vincentian Business Ethics Conference*, 2018.

Geurts, P. and Louppe, G. Learning to rank with extremely randomized trees. In *Proceedings of the Learning to Rank Challenge*, pp. 49–61, 2011.

Isayev, O., Oses, C., Toher, C., Gossett, E., Curtarolo, S., and Tropsha, A. Universal fragment descriptors for predicting properties of inorganic crystals. *Nature communications*, 8:15679, 2017.

Lemmens, A. and Croux, C. Bagging and boosting classification trees to predict churn. *Journal of Marketing Research*, 43(2):276–286, 2006.

Lin, X., Li, C., Zhang, Y., Su, B., Fan, M., and Wei, H. Selecting feature subsets based on svm-rfe and the overlapping ratio with applications in bioinformatics. *Molecules*, 23(1):52, 2017.

López, J., Maldonado, S., and Carrasco, M. Double regularization methods for robust feature selection and svm classification via dc programming. *Information Sciences*, 429:377–389, 2018.

Louppe, G. Bayesian optimisation with scikit-optimize. 2017.

Peter, S., Diego, F., Hamprecht, F. A., and Nadler, B. Cost efficient gradient boosting. In *Advances in Neural Information Processing Systems*, pp. 1551–1561, 2017.

Pinto, A., Pereira, S., Rasteiro, D., and Silva, C. A. Hierarchical brain tumour segmentation using extremely randomized trees. *Pattern Recognition*, 82:105–117, 2018.

Rahman, M. S., Rahman, M. K., Kaykobad, M., and Rahman, M. S. isgpt: An optimized model to identify sub-golgi protein types using svm and random forest based feature selection. *Artificial intelligence in medicine*, 84: 90–100, 2018.

Soltaninejad, M., Yang, G., Lambrou, T., Allinson, N., Jones, T. L., Barrick, T. R., Howe, F. A., and Ye, X.

Automated brain tumour detection and segmentation using superpixel-based extremely randomized trees in flair mri. *International journal of computer assisted radiology and surgery*, 12(2):183–203, 2017.

Wei, Z.-S., Han, K., Yang, J.-Y., Shen, H.-B., and Yu, D.-J. Protein–protein interaction sites prediction by ensembling svm and sample-weighted random forests. *Neurocomputing*, 193:201–212, 2016.

Xiao, J.-x., Lu, Z.-c., and Xu, Q.-h. A new android malicious application detection method using feature importance score. In *Proceedings of the 2018 2nd International Conference on Computer Science and Artificial Intelligence*, pp. 145–150. ACM, 2018.

A. Supplementary Material

Table 1. Descriptive statistics of the dataset.

| FEATURE	N	μ	σ	max	min
ETHICALITY	334	3.936	0.61	1.7	5
PRESSURE	334	2.700	0.78	1	5
AGE: 18-29	334	0.994	0.08	0	1
AGE: 30-49	334	0.003	0.06	0	1
AGE: 65+	334	0.003	0.06	0	1
SEX	334	0.554	0.50	0	1
M: NVR MARRIED	334	0.976	0.15	0	1
M: NOW MARRIED	334	0.018	0.13	0	1
M: SEPARATED	334	0.006	0.08	0	1
LVL EDUCATION	334	2.919	0.73	2	6
E: EMPLOYED	334	0.069	0.25	0	1
E: OUT OF WORK	334	0.006	0.08	0	1
E: SELF EMPLOYED	334	0.018	0.13	0	1
E: STUDENT	334	0.907	0.29	0	1
RELIGIOSITY	334	2.054	0.92	1	5
INCOME	334	5.180	1.26	1	6

Figure 2. 3D partial dependence plot of *Pressure* and *Income*.